

# Introduction to probabilistic programming

David Duvenaud and James Lloyd

University of Cambridge

## **Thanks to**

Daniel M Roy (Cambridge)

Roger Grosse (MIT)

# HOW TO WRITE A BAYESIAN MODELING PAPER

1. Write down a generative model in an afternoon
2. Get 2 grad students to implement inference for a month
3. Use technical details of inference to pad half of the paper

# CAN WE DO BETTER?

## Example: Graphical Models

### Application Papers

1. Write down a graphical model
2. Perform inference using general-purpose software
3. Apply to some new problem

### Inference papers

1. Identify common structures in graphical models (e.g. chains)
2. Develop efficient inference method
3. Implement in a general-purpose software package

Modeling and inference have been disentangled

## Not all models are graphical models

What is the largest class of models available?

## Probabilistic Programs

1. A probabilistic program (PP) is any program that can depend on random choices. Can be written in any language that has a (P)RNG.
2. You can specify any (computable) prior by simply writing down a PP that generates samples
3. Any PP implicitly defines a distribution over execution traces

TODO: show a program and a histogram of its output

# PROBABILISTIC PROGRAMS VS PROBABILISTIC PROGRAMMING

## Once we've defined a prior, what do we want to do with it?

The PP defines  $P(D, N, H)$ , we choose  $D$  to be the subset of variables we observe,  $H$  the set of variables we're interested in, and  $N$  the set of variables that we're not interested in, so we'll integrate them out. We want to get to  $P(H|D)$

## Probabilistic Programming

1. Usually refers to doing inference when a PP specifies your prior.
- 2.

TODO: Show the two possibilities of conditioning in the previous program

# CAN WE DEVELOP GENERIC INFERENCE FOR ALL PPs?

Yes - rejection sampling. But can we be more efficient whilst being generic?

Yes. MCMC over execution traces.

Following Wingate et alia we represent the unconditioned PP as a parameterless function  $f$

Evaluating  $f$  results in random choices which are denoted as

$$x_k = f_k | x_1, \dots, x_{k-1} \sim p_{t_k}(\cdot | \theta_k, x_1, \dots, x_{k-1}).$$

The density / probability of a particular evaluation is then

$$p(x) = \prod_{k=1}^K p_{t_k}(x_k | \theta_k, x_1, \dots, x_{k-1}).$$

We then perform MCMC over the  $x_k$  i.e. the execution trace.

# MCMC OVER EXECUTION TRACES

1. Select a random  $x_k = f_k$  in the execution trace
2. Propose a new value  $x'_k \sim K_{t_k}(\cdot | x_k, \theta_k)$
3. Run the program to determine all subsequent choices ( $x'_l : l > k$ ), reusing current choices where possible
4. Propose moving from the state  $(x_1, \dots, x_k)$  to  $(x_1, \dots, x_{k-1}, x'_k, \dots, x'_{K'})$
5. Accept the change with the appropriate reversible jump MCMC acceptance probability, this includes terms like
  - 5.1  $K_{t_k}(x'_k | x_k, \theta_k), K_{t_k}(x_k | x'_k, \theta_k), p_{t_k}(x_k | \theta_k, x_1, \dots, x_{k-1})$
  - 5.2  $\prod_{i=k}^K p_{t_i}(x_i | \theta_i, x_1, \dots, x_{i-1}), \prod_{i=k}^{K'} p_{t'_i}(x'_i | \theta'_i, x_1, \dots, x_{k-1}, x'_k, \dots, x'_{i-1})$
  - 5.3 i.e.  $\frac{K_{t_k}(x_k | x'_k, \theta_k) \prod_{i=k}^{K'} p_{t'_i}(x'_i | \theta'_i, x_1, \dots, x_{k-1}, x'_k, \dots, x'_{i-1})}{K_{t_k}(x'_k | x_k, \theta_k) \prod_{i=k}^K p_{t_i}(x_i | \theta_i, x_1, \dots, x_{i-1})}$



# FURTHER GENERAL INFERENCE METHODS

e.g. HMC, parallel tempering, etc. Remember graphical models (fancy algorithms that work in certain model classes)

# PP TIMELINE

Infer.net?

# EXAMPLE: MIXTURE OF GAUSSIANS

-

## Generative model

$$(\mu_i)_{i=1\dots k} \sim_{\text{iid}} \mathcal{N}(0, 1)$$

$$(\pi_i)_{i=1\dots k} \sim \text{Dir}(\alpha)$$

$$\Theta := \sum_{i=1}^k \pi_i \delta_{\mu_i}$$

$$(\theta_i)_{i=1\dots n} \sim_{\text{iid}} \Theta$$

$$(x_i)_{i=1\dots n} \sim_{\text{iid}} \mathcal{N}(\theta_i, 1)$$

## (Pseudo) MATLAB code

```
mu = randn(k,1);  
pi = dirichlet(k, alpha);  
  
for i = 1:n  
    theta = mu(mnrnd(1,pi));  
    x(i) = theta + randn;  
end
```

# EXAMPLE: INFINITE MIXTURE OF GAUSSIANS

## Change to generative model

$$\Theta := \sum_{i=1}^k \pi_i \delta_{\mu_i} \rightarrow \Theta \sim \text{DP}(\alpha, \mathcal{N}(0, 1))$$

## (Pseudo) MATLAB code - stick breaking construction

```
sticks = []; atoms = [];  
for i = 1:n  
    p = rand;  
    while p > sum(sticks)  
        sticks(end+1) = (1-sum(sticks)) * betarnd(1, alpha);  
        atoms(end+1) = randn;  
    end  
    theta(i) = atoms(find(cumsum(sticks) >= p, 1, 'first'));  
end  
x = theta' + randn(n, 1);
```

1. The stick breaking construction can be applied to any base measure
2. Church provides the function  $\text{DP}_{\text{mem}}$  that takes any base measure sampling function and returns a function that samples from a sample from the corresponding Dirichlet process
3. This allows easy specification of many nonparametric models e.g. HDP based models